


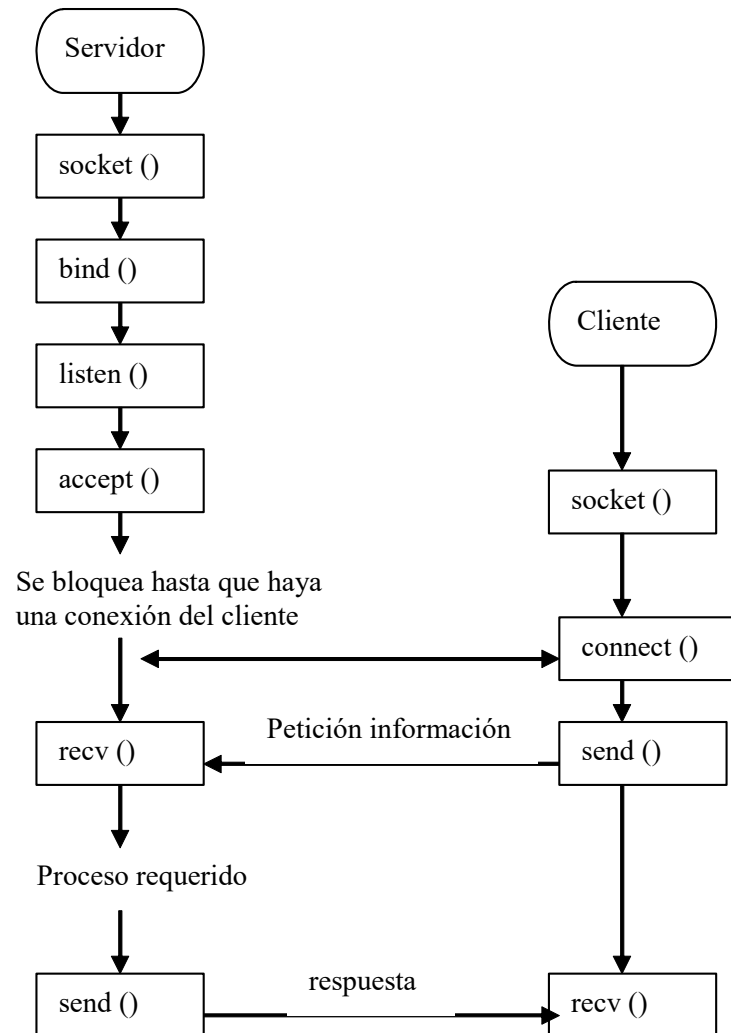
API DE PROGRAMACIÓN EN RED (SOCKETS)

Contenido


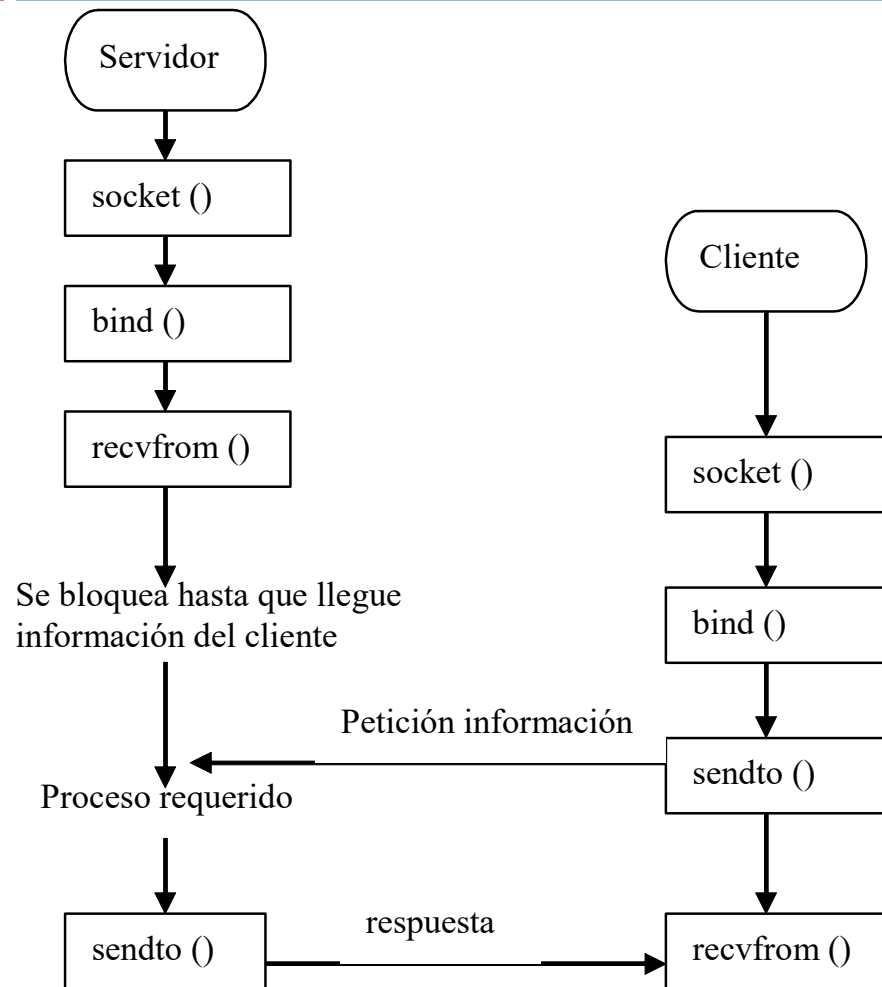
2

- Los sockets TCP y UDP
- Programas de ejemplo
 - ▣ Servidor TCP y UDP
 - ▣ Cliente TCP
 - ▣ Cliente UDP
 - ▣ Consideraciones

Los sockets de Berkeley: TCP (*stream*)



Los sockets de Berkeley: UDP (*dgram*)



Programas de ejemplo

□ servidor.c

- ▣ En TCP devuelve lo que le envía el cliente. Concretamente 1, 2, 3, 4 y 5
- ▣ En UDP devuelve la IP por la que pregunta el cliente

□ clientcp.c

- ▣ Envía 1, 2, 3, 4 y 5 y recibe lo mismo
- ▣ `clientcp nombre_servidor`

□ clientudp.c

- ▣ Pregunta al servidor por la IP dado un nombre que lee como segundo parámetro
 - `clientudp nombre_servidor nombre_cualquier_equipo`
 - Ej. `clientudp nogal www.marca.com`



TCP

Ejemplos del servidor y cliente TCP
(servidor.c y clienttcp.c)

Servidor TCP (I): *servidor.c*

7

- Crear el socket de TCP (ls_TCP) – **Socket** ()
 - ▣ Familia AF_INET (Sockets de Internet con IPv4)
 - ▣ Protocolo de transporte de tipo flujo o corriente de datos
 - SOCK_STREAM
 - ▣ Protocolo de transporte

```
50 int s_TCP, s_UDP;          /* connected socket descriptor */
93 ls_TCP = socket (AF_INET, SOCK_STREAM, 0);
94 if (ls_TCP == -1) {
95     perror(argv[0]);
96     fprintf(stderr, "%s: unable to create socket TCP\n", argv[0]);
97     exit(1);
98 }
```

Servidor TCP (II): servidor.c

- Inicializar una estructura de tipo `sockaddr_in` con la información del socket
 - ▣ Familia de direcciones
 - ▣ Dirección IP (la(s) Dir. IP propia)
 - ▣ Número de puerto “bien conocido por los clientes”

- ▣ *Dos posibles ordenaciones de los octetos de tipo entero:*
 - *Primer octeto es el más significativo (BigEndian). Por ejemplo Motorola*
 - *Primer octeto es el menos significativo (LittleEndian). Por ejemplo Intel*
- ▣ *Todo lo que se envíe por la red debe seguir la “ordenación de la red” (“BigEndian”).*
- ▣ *Métodos de conversión (#include <netinet/in.h>):*
 - *htons(): “Host to Network Short” (short de máquina a short de red).*
 - *htonl(): “Host to Network Long” (long de máquina a long de red).*
 - *ntohs(): “Network to Host Short” (short de la red a short de la máquina).*
 - *ntohl(): “Network to Host Long” (long de red a long de máquina)*

Servidor TCP (II): servidor.c

```
20 #define PUERTO 17278

57 struct sockaddr_in myaddr_in;    /* for local socket address */

69 memset ((char *)&myaddr_in, 0, sizeof(struct sockaddr_in));

85 myaddr_in.sin_family = AF_INET;
86 /* The server should listen on the wildcard address,
87  * rather than its own internet address. This is
88  * generally good practice for servers, because on
89  * systems which are connected to more than one
90  * network at once will be able to have one server
91  * listening on all networks at once. Even when the
92  * host is connected to only one network, this is good
93  * practice, because it makes the server program more
94  * portable.
95  */
96 myaddr_in.sin_addr.s_addr = htonl(INADDR_ANY);
97 myaddr_in.sin_port = htons(PUERTO);
```

Servidor TCP (III): *servidor.c*

10

- Asociar (**bind**) el socket con la información previamente indicada en la estructura `sockaddr_in`
- Reservar (**listen**) una cola para guardar las peticiones pendientes de aceptación

```
97 if (bind(ls_TCP, (const struct sockaddr *) &myaddr_in, sizeof(struct sockaddr_in)) == -1) {
98     perror(argv[0]);
99     fprintf(stderr, "%s: unable to bind address TCP\n", argv[0]);
100     exit(1);
101 }
102 /* Initiate the listen on the socket so remote users
103  * can connect. The listen backlog is set to 5, which
104  * is the largest currently supported.
105  */
106 if (listen(ls_TCP, 5) == -1) {
107     perror(argv[0]);
108     fprintf(stderr, "%s: unable to listen on socket\n", argv[0]);
109     exit(1);
110 }
```

Servidor (IV): servidor.c

11

- Convertir el servidor en un proceso demonio (*daemon*)
 - ▣ Desvincular el proceso del terminal abierto (*setpgrp*)
 - ▣ Crear el proceso que hará las funciones de servidor (*fork*)
 - ▣ Ignorar la señal *SIGCHLD* (*sigaction*) para evitar procesos zombi al morir el proceso padre
 - ▣ Registrar *SIGTERM* para la finalización ordenada del programa servidor
 - ▣ Bucle infinito para que el proceso esté siempre ejecutándose

```
130
131 /* Now, all the initialization of the server is
132    * complete, and any user errors will have already
133    * been detected. Now we can fork the daemon and
134    * return to the user. We need to do a setpgrp
135    * so that the daemon will no longer be associated
136    * with the user's control terminal. This is done
137    * before the fork, so that the child will not be
138    * a process group leader. Otherwise, if the child
139    * were to open a terminal, it would become associated
140    * with that terminal as its control terminal. It is
141    * always best for the parent to do the setpgrp.
142    */
143
144 setpgrp();
145
146 switch (fork()) {
147     case -1: /* Unable to fork, for some reason. */
148         perror(argv[0]);
149         fprintf(stderr, "%s: unable to fork daemon\n", argv[0]);
150         exit(1);
151
152     case 0: /* The child process (daemon) comes here. */
153
154         /* Close stdin and stderr so that they will not
155            * be kept open. Stdout is assumed to have been
156            * redirected to some logging file, or /dev/null.
157            * From now on, the daemon will not report any
158            * error messages. This daemon will loop forever,
159            * waiting for connections and forking a child
160            * server to handle each one.
161            */
162         close(stdin);
163         close(stderr);
164
165         /* Set SIGCLD to SIG_IGN, in order to prevent
166            * the accumulation of zombies as each child
167            * terminates. This means the daemon does not
168            * have to make wait calls to clean them up.
169            */
170         if (sigaction(SIGCHLD, &sa, NULL) == -1) {
171             perror("sigaction(SIGCHLD)");
172             fprintf(stderr, "%s: unable to register the SIGCHLD signal\n", argv[0]);
173             exit(1);
174         }
175
176         /* Registrar SIGTERM para la finalización ordenada del programa servidor */
177         vec.sa_handler = (void *) finalizar;
178         vec.sa_flags = 0;
179         if (sigaction(SIGTERM, &vec, (struct sigaction *) 0) == -1) {
180             perror("sigaction(SIGTERM)");
181             public void __cdecl perror (const char *_ErrMsg)M signal\n", argv[0]);
182             exit(1);
183         }
184
185     while (!FIN) {
```

Servidor TCP (V): *servidor.c*

12

- Aceptar peticiones
 - ▣ La función **accept** devuelve un nuevo socket (`s_TCP`) a través del cual se desarrollará el diálogo con el cliente (multiplexación de conexiones en el mismo número de puerto)
 - ▣ La dirección del cliente (IP + puerto efímero) se almacena en una nueva estructura de tipo `sockaddr_in`
 - ▣ Para cada cliente que llega se crea un proceso hijo que lo atiende
 - El servidor queda liberado para aceptar nuevos clientes
- El sockets, flujo de datos (*stream*) o asociación queda identificado por: TCP, IP servidor, puerto del servidor, IP cliente, puerto del cliente.

```
216 s_TCP = accept(ls_TCP, (struct sockaddr *) &clientaddr_in, &addrlen);
217 if (s_TCP == -1) exit(1);
218 switch (fork()) {
219     case -1: /* Can't fork, just exit. */
220         exit(1);
221     case 0: /* Child process comes here. */
222         close(ls_TCP); /* Close the listen socket inherited from the daemon. */
223         serverTCP(s_TCP, clientaddr_in);
224         exit(0);
225     default: /* Daemon process comes here. */
226         /* The daemon needs to remember
227          * to close the new accept socket
228          * after forking the child. This
229          * prevents the daemon from running
230          * out of file descriptor space. It
231          * also means that when the server
232          * closes the socket, that it will
233          * allow the socket to be destroyed
234          * since it will be the last close.
235          */
236         close(s_TCP);
237 }
```

Servidor TCP (VI): servidor.c

13

□ La función serverTCP (I)

- ▣ Dada la dirección IP del cliente obtiene su nombre (getnameinfo)
 - En caso de que no sea posible se transforma en el formato decimal punto (inet_ntop)
- ▣ Se muestra la dirección IP del cliente (hostname), el número de puerto del cliente (ahora convertido al orden del host, ntohs) y la hora de llegada

```
286 void serverTCP(int s, struct sockaddr_in clientaddr_in)
287 {
288     int reqcnt = 0; /* keeps count of number of requests */
289     char buf[TAM_BUFFER]; /* This example uses TAM_BUFFER byte messages. */
290     char hostname[MAXHOST]; /* remote host's name string */
291
292     int len, len1, status;
293     struct hostent *hp; /* pointer to host info for remote host */
294     long timevar; /* contains time returned by time() */
295
296     struct linger linger; /* allow a lingering, graceful close; */
297                          /* used when setting SO_LINGER */
298
299     /* Look up the host information for the remote host
300      * that we have connected with. Its internet address
301      * was returned by the accept call, in the main
302      * daemon loop above.
303      */
304
305     status = getnameinfo((struct sockaddr *)&clientaddr_in, sizeof(clientaddr_in),
306                          hostname, MAXHOST, NULL, 0, 0);
307
308     if(status){
309         /* The information is unavailable for the remote
310          * host. Just format its internet address to be
311          * printed out in the logging information. The
312          * address will be shown in "internet dot format".
313          */
314         /* inet_ntop para interoperatividad con IPv6 */
315         if (inet_ntop(AF_INET, &(clientaddr_in.sin_addr), hostname, MAXHOST) == NULL)
316             perror(" inet_ntop \n");
317     }
318
319     /* Log a startup message. */
320     time (&timevar);
321     /* The port number must be converted first to host byte
322      * order before printing. On most hosts, this is not
323      * necessary, but the ntohs() call is included here so
324      * that this program could easily be ported to a host
325      * that does require it.
326      */
327     printf("Startup from %s port %u at %s",
328            hostname, ntohs(clientaddr_in.sin_port), (char *) ctime(&timevar));
```

Servidor TCP (VII): *servidor.c*

14

- La función serverTCP (II)
 - ▣ Configura el socket para un cierre ordenado (setsockopt)

```
333 |     linger.l_onoff  =1;  
334 |     linger.l_linger =1;  
335 |     if (setsockopt(s, SOL_SOCKET, SO_LINGER, &linger, sizeof(linger)) == -1) {  
336 |         errout(hostname);  
337 |     }
```

Servidor TCP (VIII): servidor.c

15

- La función serverTCP (III)
 - ▣ Comienza el diálogo recibiendo datos del cliente (**recv**) a través del socket (s)
 - Los datos son almacenados en una cadena de caracteres (buf)
 - La función devuelve el número de bytes recibidos
 - ▣ La sentencia sleep(1) representa las tareas que tuviera que hacer el servidor
 - ▣ Envía datos (**send**) a través del socket (s)
 - La variable buf representa los datos enviados
 - La función devuelve el número de bytes enviados

```
361 while (len = recv(s, buf, TAM_BUFFER, 0)) {
362     if (len == -1) errout(hostname); /* error from recv */
363     /* The reason this while loop exists is that there
364      * is a remote possibility of the above recv returning
365      * less than TAM_BUFFER bytes. This is because a recv returns
366      * as soon as there is some data, and will not wait for
367      * all of the requested data to arrive. Since TAM_BUFFER bytes
368      * is relatively small compared to the allowed TCP
369      * packet sizes, a partial receive is unlikely. If
370      * this example had used 2048 bytes requests instead,
371      * a partial receive would be far more likely.
372      * This loop will keep receiving until all TAM_BUFFER bytes
373      * have been received, thus guaranteeing that the
374      * next recv at the top of the loop will start at
375      * the beginning of the next request.
376      */
377     while (len < TAM_BUFFER) {
378         len1 = recv(s, &buf[len], TAM_BUFFER-len, 0);
379         if (len1 == -1) errout(hostname);
380         len += len1;
381     }
382     /* Increment the request count. */
383     reqcnt++;
384     /* This sleep simulates the processing of the
385      * request that a real server might do.
386      */
387     sleep(1);
388     /* Send a response back to the client. */
389     if (send(s, buf, TAM_BUFFER, 0) != TAM_BUFFER) errout(hostname);
390 }
```

Servidor TCP (y IX): servidor.c

16

□ La función serverTCP (IV)

- Una vez terminado el diálogo el socket (s) se cierra (close)
- Nuevamente se muestra en pantalla la dirección IP del cliente (hostname), el número de puerto del cliente (ahora convertido al orden del host, ntohs) y la hora de llegada

```
392      /* The loop has terminated, because there are no
393      * more requests to be serviced. As mentioned above,
394      * this close will block until all of the sent replies
395      * have been received by the remote host. The reason
396      * for lingering on the close is so that the server will
397      * have a better idea of when the remote has picked up
398      * all of the data. This will allow the start and finish
399      * times printed in the log file to reflect more accurately
400      * the length of time this connection was used.
401      */
402      close(s);
403
404      /* Log a finishing message. */
405      time (&timevar);
406      /* The port number must be converted first to host byte
407      * order before printing. On most hosts, this is not
408      * necessary, but the ntohs() call is included here so
409      * that this program could easily be ported to a host
410      * that does require it.
411      */
412      printf("Completed %s port %u, %d requests, at %s\n",
413            hostname, ntohs(clientaddr_in.sin_port), reqcnt, (char *) ctime(&timevar));
414  }
415 }
```


Ciente TCP (I): clienttcp.c

17

- Crear el socket (**socket**) TCP local (s)
 - ▣ SOCK_STREAM
- Inicializar la estructura `sockaddr_in` con los datos del servidor al que desea conectarse
 - ▣ Familia de direcciones
 - ▣ Dirección IP
 - El cliente recoge como argumento el nombre del servidor al que desea conectarse
 - Obtener la IP asociada al nombre (**getaddrinfo**)
 - ▣ Número de puerto bien conocido del servidor (`htons`)

```
66  /* Set up the peer address to which we will connect. */
67  servaddr_in.sin_family = AF_INET;
68  /* Get the host information for the hostname that the
69   * user passed in.
70   */
71  memset (&hints, 0, sizeof (hints));
72  hints.ai_family = AF_INET;
73  /* esta función es la recomendada para la compatibilidad con IPv6 gethostbyname queda obsoleta */
74  errcode = getaddrinfo (argv[1], NULL, &hints, &res);
75  if (errcode != 0)
76  {
77      /* Name was not found. Return a
78       * special value signifying the
79       * error.
80       */
81      fprintf(stderr, "%s: No es posible resolver la IP de %s\n",
82              argv[0], argv[1]);
83      exit(1);
84  }
85  else {
86      /* Copy address of host */
87      servaddr_in.sin_addr = ((struct sockaddr_in *) res->ai_addr)->sin_addr;
88  }
89  freeaddrinfo(res);
90  /* hp = gethostbyname (argv[1]);
91   if (hp == NULL) {
92       fprintf(stderr, "%s: %s not found in DNS system\n",
93               argv[0], argv[1]);
94       exit(1);
95   }
96   servaddr_in.sin_addr.s_addr = ((struct in_addr *) (hp->h_addr))->s_addr;
97   */
98  servaddr_in.sin_port = htons(PUERTO);
99
100  /* Create the socket. */
101  s = socket (AF_INET, SOCK_STREAM, 0);
102  if (s == -1) {
103      perror(argv[0]);
104      fprintf(stderr, "%s: unable to create socket\n", argv[0]);
105      exit(1);
106  }
```

Cliente TCP (II): clienttcp.c

18

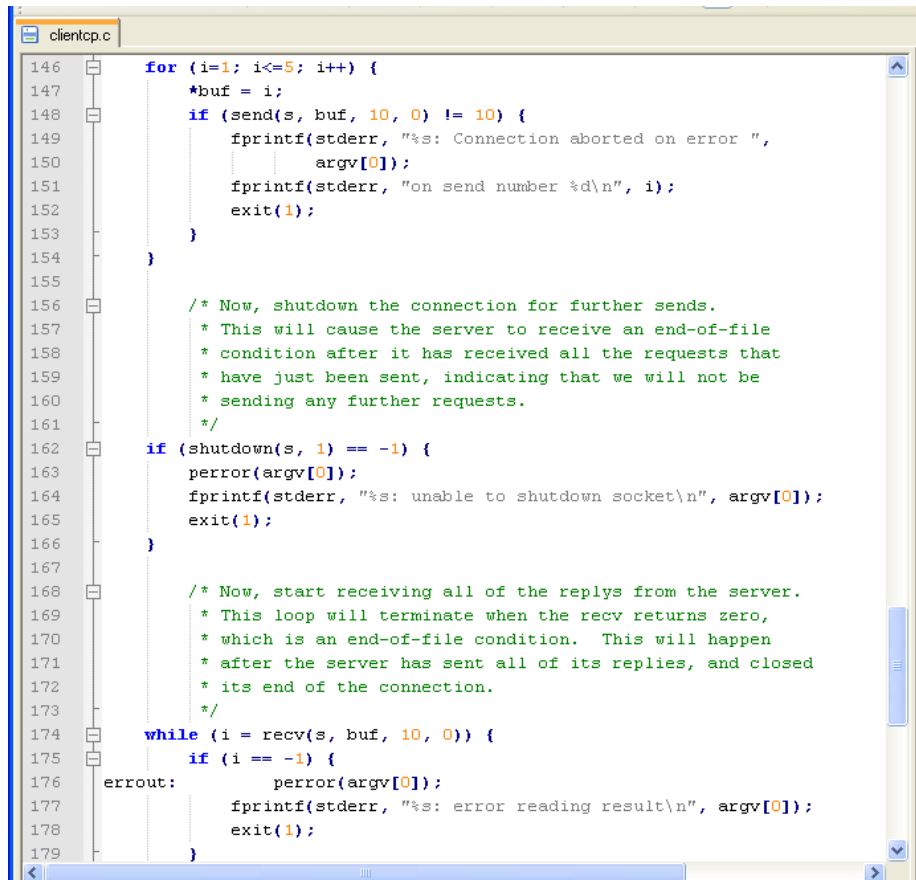
- Conectar con el servidor (**connect**)
- Si se necesita puede obtenerse la información del socket creado (**getsockname**)
 - ▣ Rellena una estructura `sockaddr_in` con la información del socket ya conectado (IP + puerto efímero)

```
clientcp.c
107      /* Try to connect to the remote server at the address
108       * which was just built into peeraddr.
109       */
110      if (connect(s, (const struct sockaddr *)&servaddr_in, sizeof(struct sockaddr_in))
111          perror(argv[0]);
112          fprintf(stderr, "%s: unable to connect to remote\n", argv[0]);
113          exit(1);
114      }
115      /* Since the connect call assigns a free address
116       * to the local end of this connection, let's use
117       * getsockname to see what it assigned. Note that
118       * addrlen needs to be passed in as a pointer,
119       * because getsockname returns the actual length
120       * of the address.
121       */
122      addrlen = sizeof(struct sockaddr_in);
123      if (getsockname(s, (struct sockaddr *)&myaddr_in, &addrlen) == -1) {
124          perror(argv[0]);
125          fprintf(stderr, "%s: unable to read socket address\n", argv[0]);
126          exit(1);
127      }
128
129      /* Print out a startup message for the user. */
130      time(&timevar);
131      /* The port number must be converted first to host byte
132       * order before printing. On most hosts, this is not
133       * necessary, but the ntohs() call is included here so
134       * that this program could easily be ported to a host
135       * that does require it.
136       */
137      printf("Connected to %s on port %u at %s",
138            argv[1], ntohs(myaddr_in.sin_port), (char *) ctime(&timevar));
139
140
```

Cliente TCP (y III): clienttcp.c

19

- Diálogo con el servidor enviando datos (**send**) y recibiendo (**recv**)
- Se puede indicar al servidor la terminación de la fase de envío de datos (**shutdown**)
- Una vez terminado el diálogo se cierra el socket (**close**)



```
clientcp.c
146 for (i=1; i<=5; i++) {
147     *buf = i;
148     if (send(s, buf, 10, 0) != 10) {
149         fprintf(stderr, "%s: Connection aborted on error ",
150             argv[0]);
151         fprintf(stderr, "on send number %d\n", i);
152         exit(1);
153     }
154 }
155
156 /* Now, shutdown the connection for further sends.
157  * This will cause the server to receive an end-of-file
158  * condition after it has received all the requests that
159  * have just been sent, indicating that we will not be
160  * sending any further requests.
161  */
162 if (shutdown(s, 1) == -1) {
163     perror(argv[0]);
164     fprintf(stderr, "%s: unable to shutdown socket\n", argv[0]);
165     exit(1);
166 }
167
168 /* Now, start receiving all of the replies from the server.
169  * This loop will terminate when the recv returns zero,
170  * which is an end-of-file condition. This will happen
171  * after the server has sent all of its replies, and closed
172  * its end of the connection.
173  */
174 while (i = recv(s, buf, 10, 0)) {
175     if (i == -1) {
176         errout:    perror(argv[0]);
177         fprintf(stderr, "%s: error reading result\n", argv[0]);
178         exit(1);
179     }
180 }
```



UDP

Ejemplos del servidor y cliente UDP
(servidor.c y clientudp.c)

Servidor UDP(I): servidor.c

21

□ Crear el socket UDP (**socket**)

- Familia de direcciones
- SOCK_DGRAM

```
116      /* Create the socket UDP. */
117      s_UDP = socket (AF_INET, SOCK_DGRAM, 0);
118      if (s_UDP == -1) {
119          perror(argv[0]);
120          printf("%s: unable to create socket UDP\n", argv[0]);
121          exit(1);
122      }
```

□ Inicializar una estructura de tipo sockaddr_in

- Familia de direcciones
- Dirección IP
- Número de puerto “bien conocido por los clientes”

```
20 #define PUERTO 17278
57      struct sockaddr_in myaddr_in;    /* for local socket address */
69      memset ((char *)&myaddr_in, 0, sizeof(struct sockaddr_in));
75      myaddr_in.sin_family = AF_INET;
76      /* The server should listen on the wildcard address,
77       * rather than its own internet address. This is
78       * generally good practice for servers, because on
79       * systems which are connected to more than one
80       * network at once will be able to have one server
81       * listening on all networks at once. Even when the
82       * host is connected to only one network, this is good
83       * practice, because it makes the server program more
84       * portable.
85       */
86      myaddr_in.sin_addr.s_addr = INADDR_ANY;
87      myaddr_in.sin_port = htons (PUERTO);
```

Servidor UDP(II): *servidor.c*

22

- Asociar (**bind**) el socket (s_UDP) con la información indicada en la estructura sockaddr_in

```
123      /* Bind the server's address to the socket. */
124      if (bind(s_UDP, (struct sockaddr *) &myaddr_in, sizeof(struct sockaddr_in)) == -1) {
125          perror(argv[0]);
126          printf("%s: unable to bind address UDP\n", argv[0]);
127          exit(1);
128      }
```

Servidor UDP(III): *servidor.c*

23

□ Atender a los clientes

- La función **recvfrom** además de recibir los datos del cliente (buffer), rellena otra estructura de tipo `sockaddr_in` (`clientaddr_in`) con los datos del cliente para poder contestarle (IP + puerto efímero)

241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260

```
/* This call will block until a new
 * request arrives. Then, it will
 * return the address of the client,
 * and a buffer containing its request.
 * BUFFERSIZE - 1 bytes are read so that
 * room is left at the end of the buffer
 * for a null character.
 */
cc = recvfrom(s_UDP, buffer, BUFFERSIZE - 1, 0,
              (struct sockaddr *)&clientaddr_in, &addrlen);
if (cc == -1) {
    perror(argv[0]);
    printf("%s: recvfrom error\n", argv[0]);
    exit (1);
}
/* Make sure the message received is
 * null terminated.
 */
buffer[cc]='\0';
serverUDP (s_UDP, buffer, clientaddr_in);
```

Servidor UDP(y IV): servudp.c

24

- `serverUDP(int s, char *buffer, struct sockaddr_in clientaddr_in)`
 - La función **getaddrinfo** dado el nombre obtiene la IP
 - Se envía (**sendto**) la dirección IP solicitada (reqaddr, una estructura de tipo in_addr)

```
454 errcode = getaddrinfo (buffer, NULL, &hints, &res);
455 if (errcode != 0)
456 {
457     /* Name was not found. Return a
458      * special value signifying the
459      * error.
460      */
461     reqaddr.s_addr = ADDRNOTFOUND;
462 }
463 else {
464     /* Copy address of host into the
465      * return buffer.
466      */
467     reqaddr = ((struct sockaddr_in *) res->ai_addr)->sin_addr;
468 }
469 freeaddrinfo(res);
470 /*
471 hp = gethostbyname (buffer);
472 if (hp == NULL) {
473     reqaddr.s_addr = ADDRNOTFOUND;
474 } else {
475     reqaddr.s_addr = ((struct in_addr *) (hp->h_addr))->s_addr;
476 }
477 */
478 /* Send the response back to the
479  * requesting client. The address
480  * is sent in network byte order. Note that
481  * all errors are ignored. The client
482  * will retry if it does not receive
483  * the response.
484  */
485 nc = sendto (s, &reqaddr, sizeof(struct in_addr),
486             0, (struct sockaddr *)&clientaddr_in, addrlen);
487 if ( nc == -1) {
488     perror("serverUDP");
489     printf("%s: sendto error\n", "serverUDP");
490     return;
491 }
492 }
```


Cliente UDP (I): clientudp.c

25

- Crear el **socket** UDP local (s)
 - ▣ Familia de direcciones
 - ▣ SOCK_DGRAM
- Inicializar servaddr_in (estructura de tipo sockaddr_in) con los datos del servidor al que desea conectarse
 - ▣ Familia de direcciones
 - ▣ Dirección IP
 - El cliente recoge como primer argumento el nombre del servidor al que desea conectarse
 - Obtiene la IP asociada al nombre (**getaddrinfo**)
 - ▣ Número de puerto del servidor (**htons**)

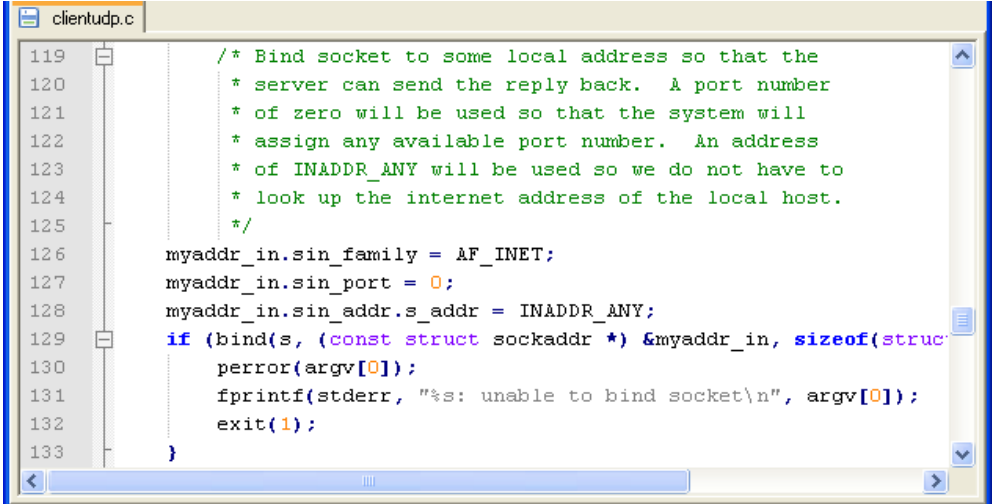
```
132 s = socket (AF_INET, SOCK_DGRAM, 0);
133 if (s == -1) {
134     perror(argv[0]);
135     fprintf(stderr, "%s: unable to create socket\n", argv[0]);
136     exit(1);
137 }
```

```
95 servaddr_in.sin_family = AF_INET;
96 /* Get the host information for the server's hostname that the
97  * user passed in.
98  */
99 memset (&hints, 0, sizeof (hints));
100 hints.ai_family = AF_INET;
101
102 /* esta función es la recomendada para la compatibilidad con IPv6 gethostbyname queda obsoleta */
103 errcode = getaddrinfo (argv[1], NULL, &hints, &res);
104 if (errcode != 0)
105 {
106     /* Name was not found. Return a
107     * special value signifying the
108     * error.
109     */
110     fprintf(stderr, "%s: No es posible resolver la IP de %s\n",
111             argv[0], argv[1]);
112     exit(1);
113 }
114 else {
115     /* Copy address of host */
116     servaddr_in.sin_addr = ((struct sockaddr_in *) res->ai_addr)->sin_addr;
117 }
118 freeaddrinfo(res);
119
120 servaddr_in.sin_port = htons(PUERTO);
```

Cliente UDP (II): clientudp.c

26

- En este caso es obligatorio el uso de la función **bind** para asociar el socket (s) con la información contenida en una estructura de tipo `sockaddr_in`
 - Familia de direcciones (AF_INET)
 - Dirección IP (INADDR_ANY)
 - Número de puerto para el cliente (un 0 significa que el sistema escoja uno libre)
- `bind` rellena la estructura con el número de puerto efímero escogido



```
119      /* Bind socket to some local address so that the
120       * server can send the reply back. A port number
121       * of zero will be used so that the system will
122       * assign any available port number. An address
123       * of INADDR_ANY will be used so we do not have to
124       * look up the internet address of the local host.
125       */
126      myaddr_in.sin_family = AF_INET;
127      myaddr_in.sin_port = 0;
128      myaddr_in.sin_addr.s_addr = INADDR_ANY;
129      if (bind(s, (const struct sockaddr *) &myaddr_in, sizeof(struct
130                perror(argv[0]));
131                fprintf(stderr, "%s: unable to bind socket\n", argv[0]));
132                exit(1);
133      }
```

Cliente UDP (y III): clientudp.c

27

- Funciones para envío y recepción son

- ▣ **sendto** y **recvfrom**

- El cliente

- ▣ Envía con sendto

- El segundo argumento son los datos que se desean enviar
 - El quinto argumento (servaddr_in) contiene los datos del servidor (IP + puerto “bien conocido”)

- Recibe con recvfrom recogiendo la dirección del socket remoto

- Hay que habilitar mecanismos de *timeout* y número de reintentos puesto que *recvfrom* es bloqueante y la respuesta puede no llegar nunca

```
while (n_retry > 0) {
    /* Send the request to the nameserver. */
    if (sendto (s, argv[2], strlen(argv[2]), 0, (struct sockaddr *)&servaddr_in,
               sizeof(struct sockaddr_in)) == -1) {
        perror(argv[0]);
        fprintf(stderr, "%s: unable to send request\n", argv[0]);
        exit(1);
    }
    /* Set up a timeout so I don't hang in case the packet
     * gets lost. After all, UDP does not guarantee
     * delivery.
     */
    alarm(TIMEOUT);
    /* Wait for the reply to come in. */

    if (recvfrom (s, &reqaddr, sizeof(struct in_addr), 0,
                  (struct sockaddr *)&servaddr_in, &addrlen) == -1) {
        if (errno == EINTR) {
            /* Alarm went off and aborted the receive.
             * Need to retry the request if we have
             * not already exceeded the retry limit.
             */
            printf("attempt %d (retries %d).\n", n_retry, RETRIES);
            n_retry--;
        }
        else {
            printf("Unable to get response from");
            exit(1);
        }
    } else {
        alarm(0);
        /* Print out response. */
        if (reqaddr.s_addr == ADDRNOTFOUND)
```

Servidor – Multiplexación de entrada/salida: servidor.c

28

□ Función **select**

- ▣ Crea el conjunto de sockets o descriptores (macros FD_ZERO y FD_SET) y se selecciona (función select) el socket que ha cambiado de estado (habitualmente por la llegada de datos)
- ▣ La detección del socket que se ha activado se realiza con la macro FD_ISSET
 - Un bloque para procesar la recepción de datos a través del socket UDP (s_UDP)
 - ServerUDP
 - Otro bloque distinto para procesar la recepción de datos a través del socket TCP (ls_TCP)
 - accept crea un nuevo socket (s)
 - serverTCP

```
if ( (numfds = select(getdtablesize(), &readmask, (fd_set *)0, (fd_set *)0, &timeout)) < 0) {
    if (errno == EINTR) {
        FIN=1;
        perror("\nselect failed\n ");
    }
}
else {

    /* Comprobamos si el socket seleccionado es el socket TCP */
    if (FD_ISSET(ls_TCP, &readmask)) {

        s_TCP = accept(ls_TCP, (struct sockaddr *) &clientaddr_in, &addrlen);
        if (s_TCP == -1) exit(1);
        switch (fork()) {
            case -1: /* Can't fork, just exit. */
                exit(1);
            case 0: /* Child process comes here. */
                close(ls_TCP); /* Close the listen socket inherited from the daemon. */
                serverTCP(s_TCP, clientaddr_in);
                exit(0);
            default: /* Daemon process comes here. */
                close(s_TCP);
        }
    } /* De TCP */

    /* Comprobamos si el socket seleccionado es el socket UDP */
    if (FD_ISSET(s_UDP, &readmask)) {

        cc = recvfrom(s_UDP, buffer, BUFFERSIZE - 1, 0,
            (struct sockaddr *)&clientaddr_in, &addrlen);
        if (cc == -1) {
            perror(argv[0]);
            printf("%s: recvfrom error\n", argv[0]);
            exit(1);
        }

        /* Make sure the message received is
        * null terminated.
        */
        buffer[cc]='\0';
        serverUDP (s_UDP, buffer, clientaddr_in);
    }
}
```

Consideraciones

29

- Las funciones `send`, `recv`, `sendto` y `recvfrom` permiten enviar y recibir datos de cualquier tipo (`void *`)
 - ▣ Se pueden enviar cadenas de caracteres, estructuras, enteros, etc.
 - ▣ Sin embargo hay que tener presente que, dado que las arquitecturas de las máquinas cliente y servidor pueden ser distintas, **para garantizar la comunicación y portabilidad del código, se recomienda usar siempre cadenas de caracteres** (o estructuras con miembros de tipo cadena)
 - ▣ Estas funciones devuelven el número de caracteres realmente leídos o enviados. Si se desea tratar lo leído como tipo cadena añadir el carácter de fin de cadena.

```
cc = recvfrom(s_UDP, buffer, BUFFERSIZE - 1, 0, (struct sockaddr *)&clientaddr_in, &addrlen);
```

```
if ( cc == -1 ) {  
    perror(argv[0]);  
    printf("%s: recvfrom error\n", argv[0]);  
    exit (1);  
}
```

```
/* Make sure the message received is null terminated */
```

```
buffer[cc]='\0';
```